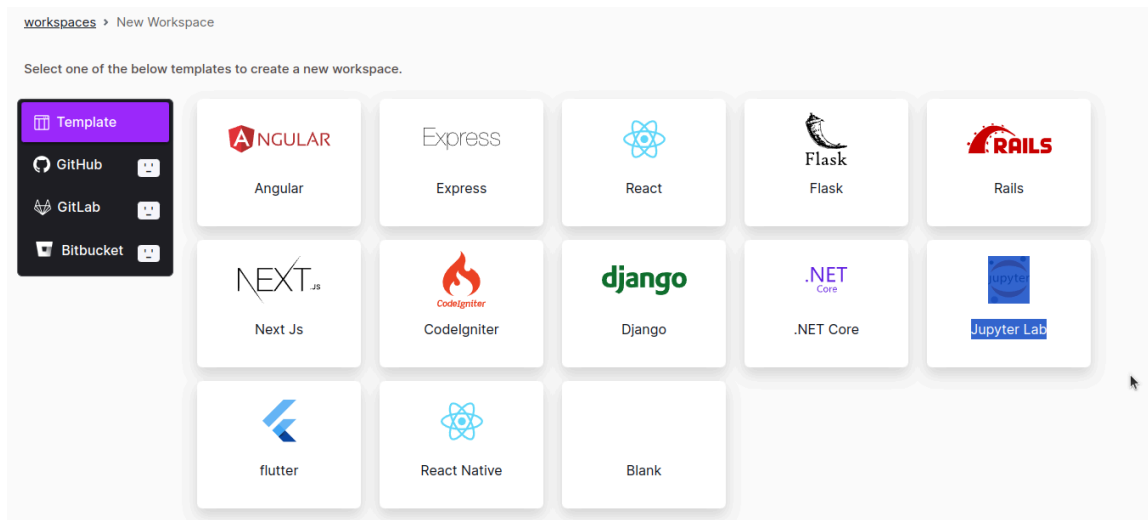
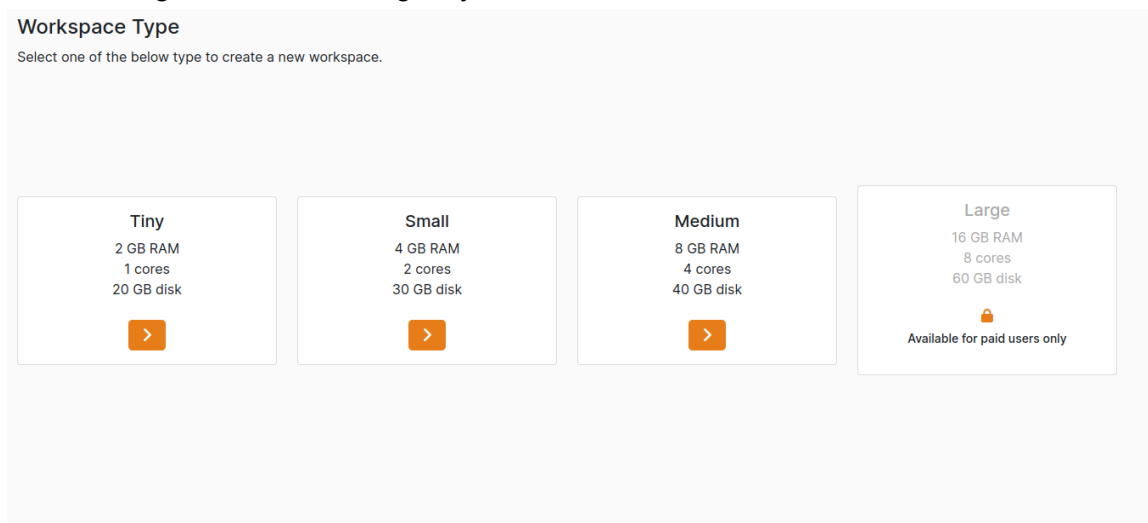


Setting up HDFS

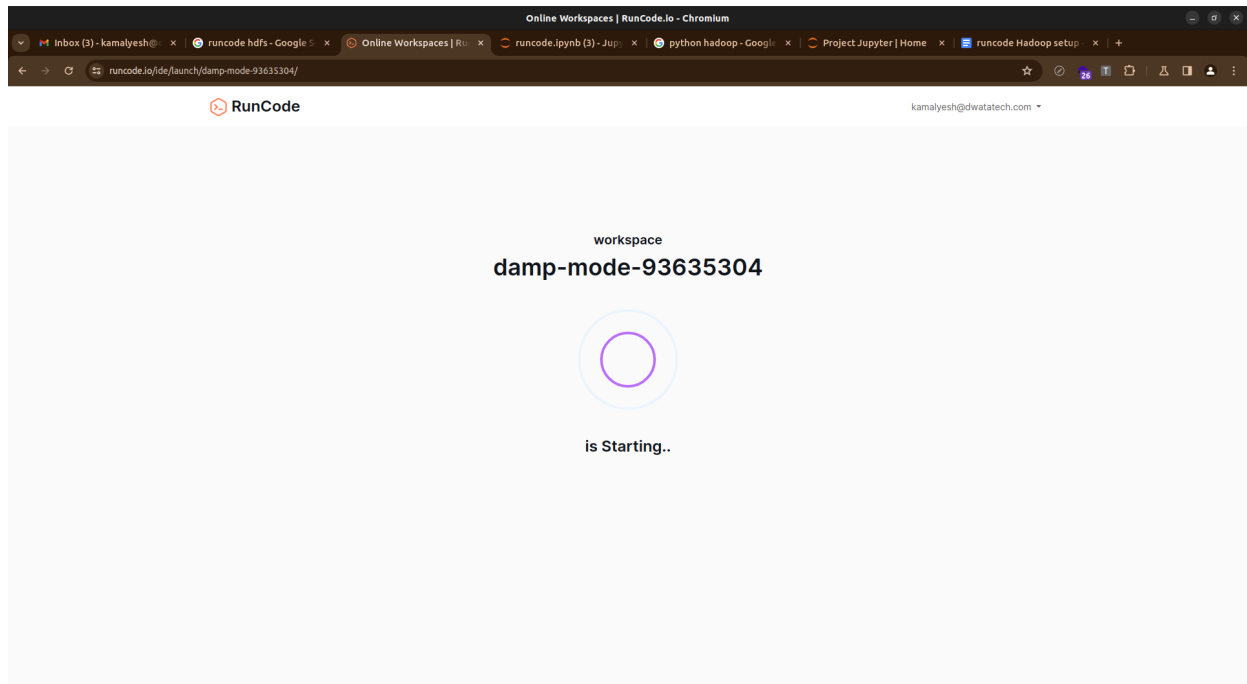
- 1 - Sign up/ Log in to runcode.io
- 2 - Create new workspace <https://runcode.io/ide/new/>



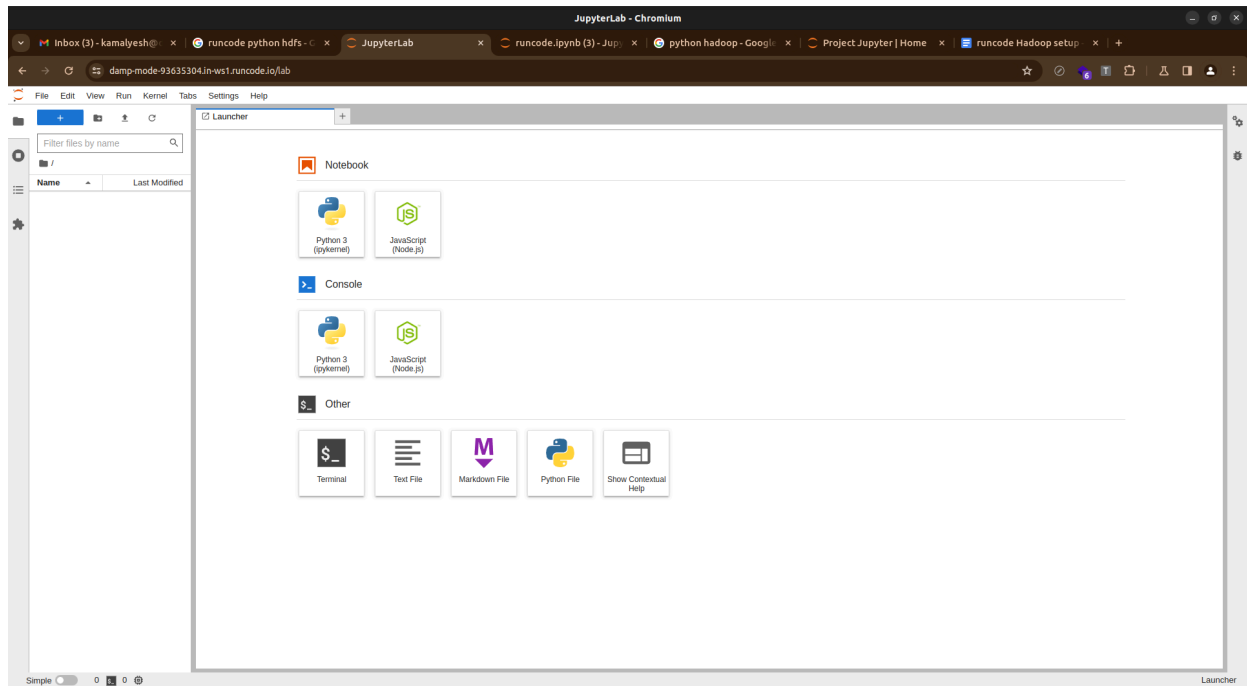
1. Select configuration. I am using Tiny for this trial



It will take some time in loading and show an error page.



Refresh the page and you will see the IDE. Open terminal



Click on the file → New → Terminal

Then start executing the below codes

For setting up HDFS follow this link - <https://sites.google.com/view/dksbin/tutorials/hadoop>

- 1 - **sudo apt-get update**
- 2 - **sudo apt install openjdk-8-jdk**
- 3 - **java -version**
- 4 - **update-alternatives --config java**

Adding a dedicated Hadoop user

- 5 - **sudo addgroup hadoop**
- 6 - **sudo adduser --ingroup hadoop hduser**
- 7 - **groups hduser**
- 8 - **sudo adduser hduser sudo**

Installing SSH

- 9 - **sudo apt-get install ssh**
- 10 - **which ssh**
- 11 - **which sshd**
- 12 - **su hduser**
- 13 - **ssh-keygen**
- 14 - **cat \$HOME/.ssh/id_rsa.pub >> \$HOME/.ssh/authorized_keys**
- 15 - **ssh localhost**

Download Hadoop

- 16 - **wget**
<https://archive.apache.org/dist/hadoop/core/hadoop-2.9.0/hadoop-2.9.0.tar.gz>
- 17 - **tar xvzf hadoop-2.9.0.tar.gz**
- 18 - **cd hadoop-2.9.0**
- 19 - **sudo mkdir -p /usr/local/hadoop**
- 20 - **sudo chown -R hduser:hadoop /usr/local/hadoop**

Setup Configuration Files

The following files need to be configured one by one

21 - `sudo nano ~/.bashrc`

Confirm the java home path according to your system

```
export JAVA_HOME=/usr/lib/jvm/java-8-openjdk-amd64
export HADOOP_INSTALL=/usr/local/hadoop
export PATH=$PATH:$HADOOP_INSTALL/bin
export PATH=$PATH:$HADOOP_INSTALL/sbin
export HADOOP_MAPRED_HOME=$HADOOP_INSTALL
export HADOOP_COMMON_HOME=$HADOOP_INSTALL
export HADOOP_HDFS_HOME=$HADOOP_INSTALL
export YARN_HOME=$HADOOP_INSTALL
export HADOOP_COMMON_LIB_NATIVE_DIR=$HADOOP_INSTALL/lib/native
export HADOOP_OPTS="-Djava.library.path=$HADOOP_INSTALL/lib"
#export HADOOP_HOME=/usr/local/hadoop/sbin
```

Add this lines to the end of the file

```
    . /usr/share/bash_completion/bash_completion
elif [ -f /etc/bash_completion ]; then
    . /etc/bash_completion
fi
fi
export JAVA_HOME=/usr/lib/jvm/java-8-openjdk-amd64
export HADOOP_INSTALL=/usr/local/hadoop
export PATH=$PATH:$HADOOP_INSTALL/bin
export PATH=$PATH:$HADOOP_INSTALL/sbin
export HADOOP_MAPRED_HOME=$HADOOP_INSTALL
export HADOOP_COMMON_HOME=$HADOOP_INSTALL
export HADOOP_HDFS_HOME=$HADOOP_INSTALL
export YARN_HOME=$HADOOP_INSTALL
export HADOOP_COMMON_LIB_NATIVE_DIR=$HADOOP_INSTALL/lib/native
export HADOOP_OPTS="-Djava.library.path=$HADOOP_INSTALL/lib"
#export HADOOP_HOME=/usr/local/hadoop/sbin
```

Like this the press Ctrl+O to overwrite and press enter, don't add the file name then press Ctrl+X to close it

22 - **source ~/.bashrc**

To save you file

23 - **javac -version**

24 - **which javac**

25 - **readlink -f /usr/bin/javac**

26 - **nano /usr/local/hadoop/etc/hadoop/hadoop-env.sh**

This is very important step to setup the java path in your hadoop environment
You have to paste this line in the above file

```
export JAVA_HOME=/usr/lib/jvm/java-8-openjdk-amd64
```

```
# Unless required by applicable law or agreed to in writing, software  
# distributed under the License is distributed on an "AS IS" BASIS,  
# WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied.  
# See the License for the specific language governing permissions and  
# limitations under the License.
```

```
# Set Hadoop-specific environment variables here.
```

```
# The only required environment variable is JAVA_HOME. All others are  
# optional. When running a distributed configuration it is best to  
# set JAVA_HOME in this file, so that it is correctly defined on  
# remote nodes.
```

```
# The java implementation to use.
```

```
export JAVA_HOME=/usr/lib/jvm/java-8-openjdk-amd64
```

```
# The jsvc implementation to use. Jsvc is required to run secure datanodes  
# that bind to privileged ports to provide authentication of data transfer  
# protocol. Jsvc is not required if SASL is configured for authentication of  
# data transfer protocol using non-privileged ports.  
#export JSVC_HOME=${JSVC_HOME}
```

```
export HADOOP_CONF_DIR=${HADOOP_CONF_DIR:-"/etc/hadoop"}
```

```
# Extra Java CLASSPATH elements. Automatically insert capacity-scheduler.  
for f in $HADOOP_HOME/contrib/capacity-scheduler/*.jar; do  
    if [ "$HADOOP_CLASSPATH" ]; then  
        export HADOOP_CLASSPATH=$HADOOP_CLASSPATH:$f  
    else  
        export HADOOP_CLASSPATH=$f  
    fi  
done
```

In the same place it will be commented just remove the hash and place your JAVA_HOME path there and

Like this the press Ctrl+O to overwrite and press enter, don't add the file name then press Ctrl+X to close it

27 - **sudo mkdir -p /app/hadoop/tmp**

28 - **sudo chown hduser:hadoop /app/hadoop/tmp**

29 - **nano /usr/local/hadoop/etc/hadoop/core-site.xml**

Add the following lines to core-site.xml

```
<configuration>
  <property>
    <name>hadoop.tmp.dir</name>
    <value>/app/hadoop/tmp</value>
    <description>A base for other temporary
directories.</description>
  </property>
  <property>
    <name>fs.default.name</name>
    <value>HDFS://localhost:54310</value>
    <description>The name of the default file system.  A URI whose
scheme and authority determine the FileSystem implementation.
The
uri's scheme determines the config property (fs.SCHEME.impl)
naming
the FileSystem implementation class.  The uri's authority is
used to
determine the host, port, etc. for a filesystem.</description>
  </property>
</configuration>
```

At the end of the file paste it </configuration> it will be written already don't paste again

Like this the press Ctrl+O to overwrite and press enter, don't add the file name then press Ctrl+X to close it

```

<!-- Put site-specific property overrides in this file. -->

<configuration>
<property>
  <name>hadoop.tmp.dir</name>
  <value>/app/hadoop/tmp</value>
  <description>A base for other temporary directories.</description>
</property>
<property>
  <name>fs.default.name</name>
  <value>hdfs://localhost:54310</value>
  <description>The name of the default file system. A URI whose
scheme and authority determine the FileSystem implementation. The
uri's scheme determines the config property (fs.SCHEME.impl) naming
the FileSystem implementation class. The uri's authority is used to
determine the host, port, etc. for a filesystem.</description>
</property>
</configuration>

```

Mapred-site.xml

30 - **cp /usr/local/hadoop/etc/hadoop/mapred-site.xml.template
/usr/local/hadoop/etc/hadoop/mapred-site.xml**

31 - **nano /usr/local/hadoop/etc/hadoop/mapred-site.xml**

Add the following lines to mapred-site.xml

```

<configuration>
  <property>
    <name>mapred.job.tracker</name>
    <value>localhost:54311</value>
    <description>The host and port that the MapReduce job tracker
runs
    at. If "local", then jobs are run in-process as a single map
    and reduce task.
    </description>
  </property>
</configuration>

```



```
-->
```

```
<!-- Put site-specific property overrides in this file. -->
```

```
<configuration>
<property>
  <name>mapred.job.tracker</name>
  <value>localhost:54311</value>
  <description>The host and port that the MapReduce job tracker runs
at. If "local", then jobs are run in-process as a single map
and reduce task.
</description>
</property>
</configuration>
```

At the end of the file paste it `</configuration>` it will be written already don't paste again

Like this the press Ctrl+O to overwrite and press enter, don't add the file name then press Ctrl+X to close it

32 - **sudo mkdir -p /usr/local/hadoop_store/HDFS/namenode**

33 - **sudo mkdir -p /usr/local/hadoop_store/HDFS/datanode**

34 - **sudo chown -R hduser:hadoop /usr/local/hadoop_store**

HDFS-site.xml

35 - **nano /usr/local/hadoop/etc/hadoop/HDFS-site.xml**

Add the following lines to HDFS-site.xml

```
<configuration>
<property>
  <name>dfs.replication</name>
  <value>1</value>
  <description>Default block replication.
  The actual number of replications can be specified when the
file is eecreated.
```

The default is used if replication is not specified in create time.

```
</description>
</property>
<property>
  <name>dfs.namenode.name.dir</name>
  <value>file:/usr/local/hadoop_store/HDFS/namenode</value>
</property>
<property>
  <name>dfs.datanode.data.dir</name>
  <value>file:/usr/local/hadoop_store/HDFS/datanode</value>
</property>
</configuration>
```

`<!-- Put site-specific property overrides in this file. -->`

```
<configuration>
<property>
  <name>dfs.replication</name>
  <value>1</value>
  <description>Default block replication.
  The actual number of replications can be specified when the file is created.
  The default is used if replication is not specified in create time.
  </description>
</property>
<property>
  <name>dfs.namenode.name.dir</name>
  <value>file:/usr/local/hadoop_store/hdfs/namenode</value>
</property>
<property>
```

Dont add this eee it should be in the same line and it was not coming thats why it is like that

At the end of the file paste it `</configuration>` it will be written already don't paste again

Like this the press Ctrl+O to overwrite and press enter, don't add the file name then press Ctrl+X to close it

Add the following lines to yarn-site.xml

36 - nano /usr/local/hadoop/etc/hadoop/yarn-site.xml

```
<configuration>
<property>
<name>yarn.nodemanager.aux-services</name>
  <value>mapreduce_shuffle</value>
</property>
<property>
  <name>yarn.nodemanager.aux-services</name>
  <value>mapreduce_shuffle</value>
</property>
</configuration>
```

Unless required by applicable law or agreed to in writing, software distributed under the License is distributed on an "AS IS" BASIS, WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied. See the License for the specific language governing permissions and limitations under the License. See accompanying LICENSE file.

-->

```
<configuration>

<property>
<name>yarn.nodemanager.aux-services</name>
  <value>mapreduce_shuffle</value>
</property>
<property>
  <name>yarn.nodemanager.aux-services</name>
  <value>mapreduce_shuffle</value>
</property>
</configuration>
```

At the end of the file paste it `</configuration>` it will be written already don't paste again

Like this the press Ctrl+O to overwrite and press enter, don't add the file name then press Ctrl+X to close it

Format the New Hadoop Filesystem

37 - hadoop namenode -format

38 - **start-dfs.sh**

39 - **start-yarn.sh**

40 - **jps**

In jps this all should run

hduser@runcode:~/hadoop-2.9.0\$ jps

19777 ResourceManager

19953 NodeManager

20198 Jps

19257 NameNode

19434 DataNode

Otherwise you HDFS is not running properly

Now your HDFS setup is done Woohoooo

We will see how to **get** and **put** file in HDFS

First will check any file is there in HDFS or not

hadoop fs -ls /

This command will show you the files present in HDFS with their permissions and location

Put command

hadoop fs -copyFromLocal /home/ubuntu/workspace/finalF.csv /

This is put command in which you have to give the file location from your hadoop cluster and at the end you have to specify the filename which you want to store in HDFS

Get command

hadoop fs -get /finalF.csv /home/ubuntu/workspace/new

This is the get command (fetching the file) after **-get** you have to give the file name which you want to extract from HDFS and **new** is the folder I have created to store the extracted file into the folder.

But

For the first time you will get the permission denied error like this

```
hduser@runcode:~/hadoop-2.9.0$ hadoop fs -get /finalF.csv
/home/ubuntu/workspace/new
24/05/08 06:15:46 WARN util.NativeCodeLoader: Unable to load native-hadoop library
for your platform... using builtin-java classes where applicable
get: /home/ubuntu/workspace/new/finalF.csv._COPYING_ (Permission denied)
```

Then just run this command

```
sudo chown hduser:hadoop -R /home/ubuntu/workspace/new
```

new is the filename which I have given in my local

And you are set with the HDFS

If any issues are there am attaching some websites links which we referred while setting up HDFS and also our logs and codes how we setup HDFS for the first time

How to download HDFS reference links

1 - <https://cloudzy.com/blog/install-hadoop-on-ubuntu/>

2 -

<https://www.digitalocean.com/community/tutorials/how-to-install-hadoop-on-ubuntu-13-10#start-hadoop>

It is the most important link

3 - <https://phoenixnap.com/kb/install-hadoop-ubuntu>

It is also one of the most imp links


4 - <https://www.guru99.com/how-to-install-hadoop.html>


For certain errors it might help


5 - <https://sites.google.com/view/dksbin/tutorials/hadoop>

From this we have got it successful and some commands which we have written some different commands it is in day 2 txt file

The bold links are the most important ones and on the top you will get the link from which this setup is done

Day 1 -  Day 1 HDFS Setup

Day 2 -  Day 2 HDFS Setup

Day 3 -  Day 3 HDFS Setup